

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn import linear_model
```

```
df = pd.read_csv('/content/drive/MyDrive/Colab Notebooks/AirQualityUCI.csv')
df
```

	Date	Time	CO(GT)	PT08.S1(CO)	NMHC(GT)	C6H6(GT)	PT08.S2(NMHC)	NOx(GT)	PT08.S3(NOx)	NO2(GT)	PT08.S4(NO2)	PT08.S5(O3)
0	10/03/2004	18:00:00	2.6	1360.0	150.0	11.9	1046.0	166.0	1056.0	113.0	169.0	1.2
1	10/03/2004	19:00:00	2.0	1292.0	112.0	9.4	955.0	103.0	1174.0	92.0	155.0	1.1
2	10/03/2004	20:00:00	2.2	1402.0	88.0	9.0	939.0	131.0	1140.0	114.0	155.0	1.1
3	10/03/2004	21:00:00	2.2	1376.0	80.0	9.2	948.0	172.0	1092.0	122.0	158.0	1.1
4	10/03/2004	22:00:00	1.6	1272.0	51.0	6.5	836.0	131.0	1205.0	116.0	149.0	1.1
...
9466	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN
9467	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN
9468	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN
9469	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN
9470	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN

9471 rows × 17 columns



```
df.dtypes
```

```
Date          object
Time          object
CO(GT)        float64
PT08.S1(CO)    float64
NMHC(GT)       float64
C6H6(GT)       float64
PT08.S2(NMHC)  float64
NOx(GT)        float64
PT08.S3(NOx)   float64
NO2(GT)        float64
PT08.S4(NO2)   float64
PT08.S5(O3)    float64
T              float64
RH             float64
AH             float64
Unnamed: 15    float64
Unnamed: 16    float64
dtype: object
```

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 9471 entries, 0 to 9470
Data columns (total 17 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Date                  9357 non-null  object
1   Time                  9357 non-null  object
2   CO(GT)                9357 non-null  float64
3   PT08.S1(CO)           9357 non-null  float64
4   NMHC(GT)              9357 non-null  float64
5   C6H6(GT)              9357 non-null  float64
6   PT08.S2(NMHC)         9357 non-null  float64
7   NOx(GT)               9357 non-null  float64
8   PT08.S3(NOx)          9357 non-null  float64
9   NO2(GT)               9357 non-null  float64
10  PT08.S4(NO2)          9357 non-null  float64
11  PT08.S5(O3)           9357 non-null  float64
12  T                     9357 non-null  float64
13  RH                    9357 non-null  float64
14  AH                    9357 non-null  float64
15  Unnamed: 15           0 non-null     float64
16  Unnamed: 16           0 non-null     float64
dtypes: float64(15), object(2)
memory usage: 1.2+ MB
```

```
#Simply drop the whole row as the missing values is less than 5% of the total data set
df.dropna(subset = ['AH'], inplace = True, axis = 0 )

#Resetting index, as we deleted some rows
df.reset_index(drop = True, inplace = True)

df.tail(10)
```

	Date	Time	CO(GT)	PT08.S1(CO)	NMHC(GT)	C6H6(GT)	PT08.S2(NMHC)	NOx(GT)	PT08.S3(NOx)	NO2(GT)	PT08.S4(NO2)
9347	04/04/2005	05:00:00	0.5	888.0	-200.0	1.3	528.0	77.0	1077.0	53.0	98.0
9348	04/04/2005	06:00:00	1.1	1031.0	-200.0	4.4	730.0	182.0	760.0	93.0	112.0
9349	04/04/2005	07:00:00	4.0	1384.0	-200.0	17.4	1221.0	594.0	470.0	155.0	160.0
9350	04/04/2005	08:00:00	5.0	1446.0	-200.0	22.4	1362.0	586.0	415.0	174.0	177.0
9351	04/04/2005	09:00:00	3.9	1297.0	-200.0	13.6	1102.0	523.0	507.0	187.0	137.0
9352	04/04/2005	10:00:00	3.1	1314.0	-200.0	13.5	1101.0	472.0	539.0	190.0	137.0
9353	04/04/2005	11:00:00	2.4	1163.0	-200.0	11.4	1027.0	353.0	604.0	179.0	126.0
9354	04/04/2005	12:00:00	2.4	1142.0	-200.0	12.4	1063.0	293.0	603.0	175.0	124.0
9355	04/04/2005	13:00:00	2.1	1003.0	-200.0	9.5	961.0	235.0	702.0	156.0	104.0
9356	04/04/2005	14:00:00	2.2	1071.0	-200.0	11.9	1047.0	265.0	654.0	168.0	112.0



```
df.drop(['Unnamed: 15', 'Unnamed: 16'], axis = 1, inplace=True)
```

```
df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 9357 entries, 0 to 9356
Data columns (total 15 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Date                   9357 non-null   object
1   Time                   9357 non-null   object
2   CO(GT)                 9357 non-null   float64
3   PT08.S1(CO)            9357 non-null   float64
4   NMHC(GT)               9357 non-null   float64
5   C6H6(GT)               9357 non-null   float64
6   PT08.S2(NMHC)          9357 non-null   float64
7   NOx(GT)                9357 non-null   float64
8   PT08.S3(NOx)           9357 non-null   float64
9   NO2(GT)                9357 non-null   float64
10  PT08.S4(NO2)           9357 non-null   float64
11  PT08.S5(O3)            9357 non-null   float64
12  T                       9357 non-null   float64
13  RH                     9357 non-null   float64
14  AH                     9357 non-null   float64
dtypes: float64(13), object(2)
memory usage: 1.1+ MB
```

```
df['Date'] = df['Date'].astype('category')
df['Date'] = df['Date'].cat.codes

df['Time'] = df['Time'].astype('category')
df['Time'] = df['Time'].cat.codes

df.tail(10)
```

	Date	Time	CO(GT)	PT08.S1(CO)	NMHC(GT)	C6H6(GT)	PT08.S2(NMHC)	NOx(GT)	PT08.S3(NOx)	NO2(GT)	PT08.S4(NO2)	PT08
9347	43	5	0.5	888.0	-200.0	1.3	528.0	77.0	1077.0	53.0	987.0	
9348	43	6	1.1	1031.0	-200.0	4.4	730.0	182.0	760.0	93.0	1129.0	
9349	43	7	4.0	1384.0	-200.0	17.4	1221.0	594.0	470.0	155.0	1600.0	

```
df.isnull().sum()

Date      0
Time      0
CO(GT)    0
PT08.S1(CO) 0
NMHC(GT)  0
C6H6(GT)  0
PT08.S2(NMHC) 0
NOx(GT)   0
PT08.S3(NOx) 0
NO2(GT)   0
PT08.S4(NO2) 0
PT08.S5(O3) 0
T         0
RH        0
AH        0
dtype: int64

X = df.drop(columns = 'AH')
X
```

	Date	Time	CO(GT)	PT08.S1(CO)	NMHC(GT)	C6H6(GT)	PT08.S2(NMHC)	NOx(GT)	PT08.S3(NOx)	NO2(GT)	PT08.S4(NO2)	PT08
0	114	18	2.6	1360.0	150.0	11.9	1046.0	166.0	1056.0	113.0	1692.0	
1	114	19	2.0	1292.0	112.0	9.4	955.0	103.0	1174.0	92.0	1559.0	
2	114	20	2.2	1402.0	88.0	9.0	939.0	131.0	1140.0	114.0	1555.0	
3	114	21	2.2	1376.0	80.0	9.2	948.0	172.0	1092.0	122.0	1584.0	
4	114	22	1.6	1272.0	51.0	6.5	836.0	131.0	1205.0	116.0	1490.0	
...	
9352	43	10	3.1	1314.0	-200.0	13.5	1101.0	472.0	539.0	190.0	1374.0	
9353	43	11	2.4	1163.0	-200.0	11.4	1027.0	353.0	604.0	179.0	1264.0	
9354	43	12	2.4	1142.0	-200.0	12.4	1063.0	293.0	603.0	175.0	1241.0	
9355	43	13	2.1	1003.0	-200.0	9.5	961.0	235.0	702.0	156.0	1041.0	
9356	43	14	2.2	1071.0	-200.0	11.9	1047.0	265.0	654.0	168.0	1129.0	

9357 rows × 14 columns

```
y = df['AH']

from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=0)

from sklearn.preprocessing import StandardScaler
scaler = StandardScaler()
X_train_scaler = scaler.fit_transform(X_train)
X_test_scaler = scaler.transform(X_test)

from sklearn.linear_model import LinearRegression
lr = LinearRegression()

# Fitting our training data to our linear regression model
lr.fit(X_train_scaler, y_train)

LinearRegression()

c = lr.intercept_
c

-6.953743563902885

m = lr.coef_
```


m

```
array([-5.21915711e-02,  1.52694065e-01, -5.57130177e-03, -2.76042444e-01,
        2.63557243e-01,  3.20174099e+01, -6.21895189e+00, -8.43389953e-01,
       -2.97894262e-01,  6.63583247e-01, -2.26458983e+00,  4.44067865e-01,
        9.86181865e+00,  3.95902365e+00])
```

```
y_pred_train = lr.predict(X_train_scaler)
y_pred_train
```

```
array([0.87446851, 2.32166635, 1.568725 , ..., 2.23828705, 1.74218483,
       1.86518636])
```

```
plt.scatter(y_train,y_pred_train)
plt.xlabel("Actual Absolute Humidity")
plt.ylabel("Predicted Absolute Humidity")
plt.show()
```




```
from sklearn.metrics import r2_score
r2_score(y_train, y_pred_train)

0.9993336584579668
```

```
y_pred_test = lr.predict(X_test_scaler)
```

```
plt.scatter(y_test,y_pred_test)
plt.xlabel("Actual Absolute Humidity")
plt.ylabel("Predicted Absolute Humidity")
plt.show()
```



```
r2_score(y_test, y_pred_test)
```

0.9993618893773653

Colab paid products - Cancel contracts here			
✓	0s	completed at 19:24	<div><div></div><div>×</div></div>